

COSI 131a: Fall 2015

Programming Assignment 2

Overview

Your project is to implement a simulation of customers issuing requests to use shared resources using Java Threads. There are rules that constrain when different customers may make use of resources. Your program is responsible for enforcing these rules using Java synchronization.

The project is split into two tasks. In task 1, you will have multiple concurrent client threads representing consumers that will try to access a shared resource.

To enforce the restrictions on which and how many consumers can use the resource your simulation program will need to use Java synchronized methods and busy waiting.

In task 2, you will improve the simulation program by using Java monitors to avoid busy waiting. In this task you will have a centralized scheduler that coordinates consumer access to multiple instances of the shared resources. Your scheduler will build on the synchronized code you developed in task 1, and add monitor synchronization to avoid busy waiting. The scheduler will also need to enforce FIFO (First-In First-Out) ordering.

Task 1 has an earlier due date than Task 2. See LATTE for due dates.

Consumer types and rules governing consumer access.

The following rules limit, subject to customer type, when, where and with whom client threads representing customers can make requests of resources. Your code must enforce these rules in both Task 1 and Task 2.

- Client (General)
 - Have the following customer attributes: type, name, industry, request level, speed.
 - **Type** can be *shared* or *basic*.
 - Customers of the same **industry** do not like being around other customers of the same industry so concurrent clients at the same server must be of a different **industry**.
 - **Request level** determines the number of requests that must be processed before the client ceases operations. A client with request level 0 will not attempt to connect to a server. Each successful connection decrements request level. A client starts with level 3.
 - **Speed** determines how many milliseconds it takes for a client's request to be processed at a server after successfully gaining access to said server (higher speed value = less waiting). The exact formula for waiting time is $((10 - \text{speed}) * 100) \text{ ms}$. I.e.

a client with speed 8 will gain access to a server, wait for 200ms then “leave” the server. A client can have a speed between 0 and 9. The speed can be passed in, or randomly generated.

- Every client must have a toString() method that returns “INDUSTRY TYPE NAME” i.e. “TECH BASIC AA1”.

- Servers (General)

- Clients make a request to utilize a server by calling Server.connect()
- Server.connect() will return true (client has successfully connected) if the client’s joining would not violate any of the constraints of the server (outlined under Basic Server). Otherwise Server.connect() returns false.

- Basic Server (Task 1)

- An instance of a basic server will not service:
 - more than two shared clients at the same time.
 - more than one basic client at a time.
 - a shared and basic client at the same time.
 - two shared clients with the same industry.

- Master Server (Task 2)

- Distributes clients to their respective basic servers.
- Keeps list of references to basic servers (in a HashMap).
- Keeps list of references to the queues of each basic server.

As you know from learning about concurrency, without proper synchronization, a client thread trying to enforce these constraints is prone to race conditions. E.g. a thread could check to see it can use a server, then be interrupted before it can actually use the server, at which point another thread may also check to see if it can use the server and do so. When the first thread runs again and proceeds to use the server, one of the invariants may be violated (e.g., maybe a basic client and shared client will be using the same server at the same time).

Part of your job in this assignment is to use synchronization to prevent such race conditions from occurring.

To implement your solutions you will need to use the code base provided by us, described below. This code base includes a test harness that allows us (and you) to test the correctness of your solutions. You will be therefore will be instructed to follow certain conventions in your code and will be not allowed to modify certain classes.

For example, your code will not be starting or joining client threads. Instead, the threads will be started by the test harness. Please follow the the instructions carefully.

You will need to understand how the entire provided code works to complete your assignment.

Provided Code

The code for this assignment is divided into the following packages

`edu.brandeis.cs131.Common.Abstract`

The Abstract package contains the abstract classes and interfaces that your program must implement. You are not allowed to modify any of the files within this package.

`edu.brandeis.cs131.Common.Abstract.Log`

The Abstract.Log package contains classes used to record the actions of clients. These classes are used by the test packages to evaluate that constraints have not been violated. The log is not a means for passing messages. The only reference to Log classes in your code should be in the `ConcreteFactory.createNewMasterServer` method. You are not allowed to modify any of the files within this package.

`edu.brandeis.cs131.Common.YourName`

The *YourName* package contains a few mostly empty classes that implement the classes found in the Abstract package. You should initially rename this package to your first then last name with no spaces i.e. `JaneSmith`. All of the code pertaining to your solution of this assignment must go in this package. You are free to add, edit, rename, and delete files in this package in the course of implementing your solutions.

Test Files - `edu.brandeis.cs131.Common.Test`

You must modify `TestUtilities` to include your version of `ConcreteFactory` (Modify the import statement to include your `ConcreteFactory`), in order for the test suite to run. You may not edit any of the other test classes, but we highly encourage you to study them for your understanding and to aid in discovering why a test is failing.

API Documentation

You will need to understand how all provided code works to solve this problem.

You will want to inspect the java files yourself, but to get you started we include the API docs generated with javadoc from the files we provide.

Task 1: The Basic Server

You must make class(es) that extend the abstract class Client and display the behavior described previously.

You must create a class that extends the abstract class Server and carries out the following tasks:

- Enforce the Server Restrictions described previously.
- Use the Java synchronized keyword to achieve mutual exclusion in the critical sections and prevent race conditions.

You must setup your ConcreteFactory to create and return your newly created classes.

When this task is complete your code should pass all tests included in BehaviorTest, and SimulationTest.Basic_Server_Test.

Task 2: The Master Server

You have now successfully programmed BasicServer which enforces entry criteria and prevents race conditions. However, there are two problems with the design of BasicServer:

- When the server is not available to a client, that client busy-waits (loops over the list of known servers repeatedly inside of its run() method).

In this task, you will solve the same problem but by using Java monitor synchronization to avoid busy waiting. In addition each client is designated an instance of a basic server (handled by the hashmap in the master server) based on its speed.

When this task is complete your code should pass all tests included with this assignment.

Implementation

You must create a new class (MasterServer) that extends Server and contains a map of BasicServers, each of which have their own First In First Out wait-queue. The client in the front of a BasicServer's wait-queue will attempt to connect to that server, and if unsuccessful wait for a client currently connected to that server to disconnect. BasicServer should be the same class you implemented in the first part of this project.

MasterServer will carry out the following tasks:

- Keep references to BasicServers and their queues as private member variables inside MasterServer.
- When a call is made to connect(client) on a MasterServer instance, the following general steps should be executed:

- The client should use the provided `getKey()` method to determine what server it has been assigned access to. If there is a non-empty queue for this server the client adds itself to the tail of the queue and waits for its turn. Only the head of the queue is ever allowed to proceed.
 - That client will then try to make a request to that server. If it cannot make a request due to a server condition being violated, the client thread must wait. Otherwise, the client successfully starts making a request to a server.
 - Only the head of a queue should be notified when a client leaves a server
 - When a client successfully connects to a server it should remove itself from the queue and notify the new head of the queue. So it may attempt to connect, remember two Shared Clients can enter a server at once.
- When a Client exits the MasterServer by calling `disconnect` on a MasterServer instance, the MasterServer must call `disconnect()` on the BasicServer the client is currently connected to. Remember, you may not modify `Server.java` or any of the other classes provided to you, and should not have to change your implementation of BasicServer, so you must solve this problem within MasterServer.
 - Use monitors to avoid busy waiting if the client's assigned server is busy.
 - Maximize concurrency by synchronizing only where necessary.
 - A client only removes itself from the queue, if it successfully requests from a BasicServer, otherwise it waits.
 - Only the head of the queue tries to enter a BasicServer. All other clients wait their turn (even if they could enter simultaneously).

No 'main' method

Note that your only point of entry in the code we provide is through the JUnit tests, which will set up the environment in which your client threads will run. Your tasks for this assignment do not include writing a main method. Rather, you must rely on your understanding of busy-waiting/monitors, and the JUnit tests to know if you have completed the assignment.

Important Note about Disallowed Java Tools

In PA1, you were instructed to consider using a synchronized class provided by Java for inter-thread communication (`LinkedBlockingQueue`) to solve the problem. For this project, that is not allowed; you may not use any synchronized data structure included in the Java API. You must write your own (using the "synchronized" keyword). Of course, you can and should use non-synchronized data structures in the Java standard library. You can consult the API docs to see if a data structure is synchronized.

Submission Guidelines

Create a zip or tar.gz file containing your completed code. You should name this file: last name underscore first name. For example if your name is Jane Smith you should name your file: Smith_Jane.zip or Smith_Jane.tar.gz. This file should contain the src folder within the archive not as the top level (when the archive is opened the src folder should be visible). Upload this archive to the submission folder on Latte. Please also include a README file that contains a description of your implementation and any issues your code might have. This is an individual project and all submitted work must be your own. To this effect, to receive any credit, you will have to include in your README a statement, attesting the submitted work is yours. So while you are allowed to discuss your solution with other students, make sure you do not share your code. If you have extensive discussions with another student please indicate this in your README file.

Helpful Hints

'instanceof' operator

You might find the instanceof operator (scroll down a bit on the page that links to) helpful. instanceof can tell you if an object can be downcast from a parent type into a child type. Typically, we don't use it because there are better techniques (polymorphism leverages the type system to do the work for you), but for this problem it will probably help you out.

Here is an example:

```
public class TestInstanceOf {
    private static class Parent {
    }
    private static class Child1 extends Parent {
    }
    private static class Child2 extends Parent {
    }

    public static void main(String[] args) {
        Parent p1 = new Child1();
        Parent p2 = new Child2();

        if(p1 instanceof Child1) {
            System.out.println("p1 is instance of Child1");
        }
        if(p1 instanceof Child2) {
```

```

        System.out.println("p1 is instance of Child2");
    }

    if(p2 instanceof Child1) {
        System.out.println("p2 is instance of Child1");
    }
    if(p2 instanceof Child2) {
        System.out.println("p2 is instance of Child2");
    }
}
}

```

This outputs:

```

p1 is instance of Child1
p2 is instance of Child2

```

'synchronized' blocks

The following two uses of the synchronized keyword are equivalent:

```

public synchronized void m() {
    // ...
}
public void m() {
    synchronized(this) {
        // ...
    }
}

```

If you want to synchronize the whole method on the lock contained inside of "this", then it is more elegant to put the synchronized keyword in the signature of the method. Sometimes, you want to synchronize on smaller blocks of code than an entire method. This allows you to limit the amount of synchronization, which can improve performance by maximizing concurrency. In this case, putting synchronized in its own block can be the better choice.

```

public void m() {
    // a noncritical section
    synchronized(this) {
        // a critical section
    }
    // a noncritical section
    synchronized(this) {

```

```
    // a critical section
}

// a noncritical section
}
```

Note that you can synchronize on the lock in any instance of an object, not just "this". For example:

```
public class C {
    private static Object o = new Object();

    public void m() {
        synchronized(o) {
            // ...
        }
    }
}
```

You might need to synchronize on an instance other than "this" for this project.